

Exploring the limits of Transfer learning with a Unified Text-to-Text Transformer

1. Introduction

2. Setup

2.1 Model

2.2. The Colossal Clean Crawled Corpus

2.3 Downstream Tasks

2.4 Input and Output Format

3. Experiments

3.1 Baseline

3.1.1. Model

3.1.2 Training

3.1.3 Vocabulary

3.1.4 Unsupervised Objective

3.1.5 Baseline Performance

3.2 Architectures

3.2.1 Model Structures

3.2.2 Comparing Different Model Structures

3.2.3 Objectives

3.2.4 Results

3.3 Unsupervised Objectives

3.3.1 Disparate High-level Approaches

3.3.2 Simplifying the BERT Objective

3.3.3 Varying the Corruption Rate

3.3.4 Corruption Span

3.3.5 Discussion

3.4 Pre-training Data set

3.4.1 Unlabeled Data Sets

3.4.2 Pre-training Data Set Size

3.5 Training Strategy

3.5.1 Fine-tuning Methods

3.5.2 Multi-task Learning

3.5.3 Combining Multi-Task Learning with Fine-Tuning

3.6 Scaling

3.7 Putting it All Together

4. Reflection

4.1 Takeaways

4.2 Outlook

2.1. Model

- roughly equivalent to the original Transformer.
- removing the Layer Norm bias
- placing the layer normalization outside the residual path
- using different position embedding scheme.

2.4. Input & Output Format

- "text-to-text" format
- provides a consistent training objective both for pre-training & fine-tuning

3.1. Baseline

- to reflect modern, typical practice
- pre-train a standard Transformer using a simple denoising objective
- separately fine-tune on each of downstream task.

3.1.2 Training

- "inverse square root" learning rate schedule

$$1 / \sqrt{\max(n, k)}$$

current training iteration number of warm-up steps (set to 10^4 in all experiments)

- constant learning rate of 0.01 for the first 10^4 steps
- then exponentially decays the learning rate until pre-training is over.

3.1.4. Unsupervised Objective

- mask consecutive spans of tokens & only predict dropped-out tokens
→ reduce computational cost of pre-training

3.1.5. Baseline Performance

- pre-training provides significant gains across almost all benchmarks
- inter-run variance
 - GLUE & SuperGLUE : average of scores of each benchmark
 - CoLA, CB, COPA : ↑ inter-run variance...
 - harder to compare models using GLUE & SuperGLUE scores alone.

3.2 Architectures

3.2.1. Model Structures

- major distinguishing factor for different architectures
 - : "mask" used by different attention mechanisms in the model.

- Transformer

- input: sequence
- output: a new sequence of the same length
 - each entry of output sequence is produced by computing a weighted average of entries of the input sequence

$$y_i = \sum_j w_{i,j} x_j$$

y_i : i^{th} element of the output sequence
 $w_{i,j}$: scalar weight self-attention mechanism as a function of x_i & x_j
 x_j : j^{th} entry of the input sequence

- attention mask is used to zero out certain weights in order to constrain which entries of the input can be attended to at a given output timestep.

① encoder - decoder Transformer

- encoder:

- fully-visible attention mask
- allows a self-attention mechanism to attend to any entry of the input when producing each entry of its output.
- BERT's output at the timestep corresponding to the classification token is then used to make a prediction for classifying the input sequence.

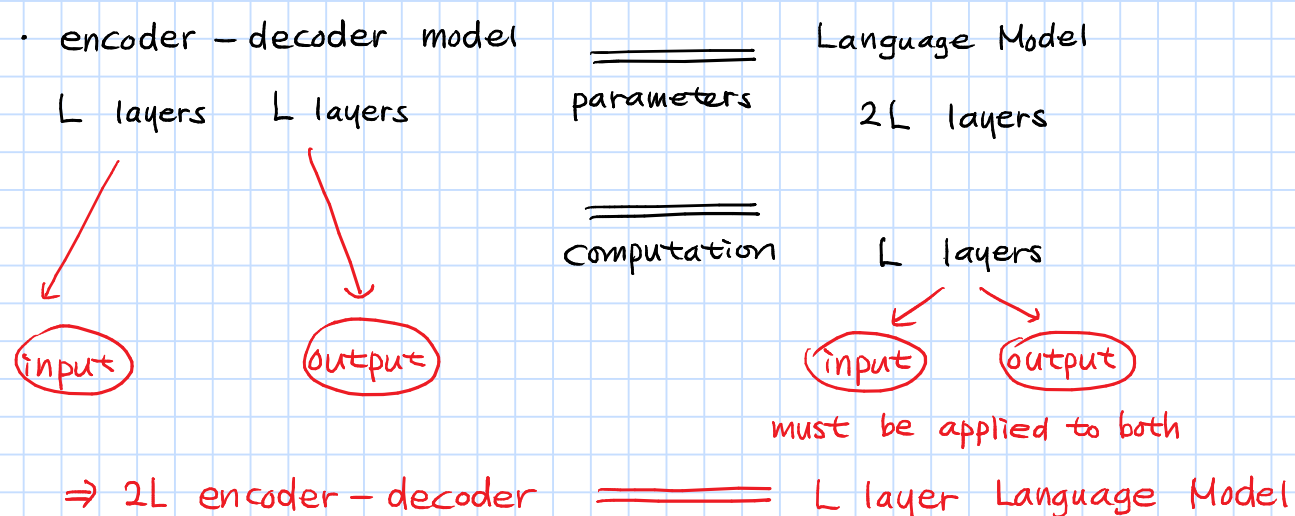
- decoder:

- causal masking pattern
- used to autoregressively produce an output sequence
- language models can learn to perform some text-to-text tasks without supervision.
- drawback of causal masking in the text-to-text setting
 - avoided by changing the masking pattern.

3.2.2. Comparing Different Model Structure

- equivalent in some meaningful way: same number of parameters, same amount of computation

- hard! why?



- Comparing different model structures

- encoder-decoder L layers: 2P, M
 - encoder-decoder shared L layers: P, M
 - encoder-decoder L/2 layers: P, M/2
 - decoder-only LM: P, M
 - decoder-only prefix LM: P, M
- parameters FLOPs

3.2.3. Objectives

- for models that ingest a prefix before making predictions:

(encoder-decoder model & prefix LM)

- sample a span of text from unlabeled data set
- choose a random point to split it into prefix and target points.

- for standard language model:

- train the model to predict the entire span from beginning to end.

3.2.4. Results

- shared parameter encoder-decoder outperforms the decoder-only prefix LM
- explicit encoder-decoder attention is beneficial.

3.3.1. Disparate High-Level Approaches

① prefix language modeling

- splits a span of text into two components: input to the encoder / target sequence predicted by the decoder
- Example:

Thank you for inviting me to your party last night.

inputs targets

② masked language modeling (MLM) objected used in BERT

- takes a span of text & corrupts 15% of the tokens
- 90% of corrupted tokens : replaced with a special mask token
- 10% " : replaced with a random token
- BERT goal: reconstruct masked tokens at the output of the encoder.
- encoder - decoder case : use the entire uncorrupted sequence as target
- Example :

Thank you <M><M> me to your party apple week. (original text)

inputs target

③ Deshuffling objective

- denoising sequential autoencoder
- takes a sequence of tokens, shuffles it, uses the original deshuffled sequence as a target
- Example :

party me for your to. last fun you inviting week Thank (original text)

inputs target

3.3.4. Corrupting Spans

- specifically corrupt contiguous, randomly-spaced spans of tokens.
- Example :

- a sequence of 500 tokens, 15% corrupted, 25 total spans

↓

- total number of corrupted tokens : $500 \times 0.15 = 75$
- average span length : $75 / 25 = 3$

→ given the original sequence length & corruption rate,
we can equivalently parametrize this objective
by the average span length or total number of spans.